

Vishrut Ghotge

Bengaluru, India | vishrutghotge3@gmail.com | +91-8668751906 | linkedin.com/in/vishrutg | github.com/Vishrut33

PROFESSIONAL SUMMARY

AI/MLOps Engineer with 1 year of experience automating enterprise systems, building CI/CD pipelines, and integrating machine learning models into production workflows. Proficient in Python, scikit-learn, XGBoost, MLflow, FastAPI, Docker, and AWS. Delivered a 40% reduction in manual testing and a 30% faster release cycle at Tata Elxsi by applying MLOps and DevOps practices on a Qualcomm client project.

TECHNICAL SKILLS

Machine Learning: scikit-learn, XGBoost, LightGBM, Pandas, NumPy, Supervised Learning, Feature Engineering, Hyperparameter Tuning, Model Evaluation

MLOps and Deployment: MLflow, FastAPI, Model Serving, Model Registry, CI/CD for ML, Model Monitoring, REST API Development

LLM and Generative AI: LangChain, Retrieval-Augmented Generation, Prompt Engineering, Ollama, Large Language Models, FAISS

Data Engineering: Apache Airflow, ETL Pipelines, SQL, PostgreSQL, Data Preprocessing

Cloud and Infrastructure: AWS (EC2, S3, VPC, IAM), Docker, Kubernetes, Terraform, Infrastructure as Code

CI/CD and Version Control: GitHub Actions, Jenkins, Git, GitLab, Automated Testing, PyTest

Monitoring and Observability: Prometheus, Grafana, Splunk

Programming: Python, Bash, Java, JavaScript

Systems: Linux, TCP/IP, DNS, Agile

EXPERIENCE

AI/MLOps Engineer

Jun 2024 – Present

Tata Elxsi, Client: Qualcomm

Bengaluru, India

- Cut manual testing effort by **40%** by engineering Python automation frameworks with PyTest, Boardfarm, and MTF5 for network-system validation across 3 product lines.
- Reduced release cycle time by **30%** by building Jenkins and GitHub Actions CI/CD pipelines that automated build, test, and deployment across 5 environments.
- Decreased mean time to resolution (MTTR) by **25%** by configuring Prometheus and Grafana dashboards tracking 15+ system health metrics with automated alerting.
- Resolved **20+** critical Linux system failures through structured log analysis and performance profiling, maintaining 99% uptime on the Qualcomm client project.
- Accelerated team onboarding by **2 weeks** by documenting automation standards and aligning frameworks with MLOps and DevOps best practices across a 6-person engineering team.

PROJECTS

ATLAS: Algorithmic Trading System with ML Signal Filtering

2024 – Present

- Reduced false trade entries by **35%** by training an XGBoost classifier on 12 OHLCV-derived features to filter low-confidence signals across 5 forex pairs on MetaTrader 5.
- Processed **10,000+** historical candles per symbol using a Pandas and NumPy feature pipeline to generate inputs for the ML signal classifier.
- Raised signal consistency by **20%** by weighting outputs from 6 technical analysis strategies through a dynamic scoring engine.
- Validated trade signals with a 5-persona Ollama LLM consensus layer, combining rule-based logic with generative AI to reduce conflicting entries.
- Maintained a modular Python codebase with structured logging, graceful error handling, and environment-level config across **3,000+** lines of production-oriented code.

MLOps Pipeline: Experiment Tracking and Model Serving

2025

- Tracked **50+** training experiments using MLflow, logging accuracy, F1-score, and AUC-ROC per run and registering the best model for one-click promotion or rollback.
- Served a production XGBoost model via FastAPI REST endpoint with Pydantic validation, containerised in Docker, achieving sub-**50ms** inference latency.

- Automated retraining with a GitHub Actions workflow that runs evaluation gates and blocks promotion if accuracy drops below a defined threshold.

RAG Document Q&A System: LangChain and FAISS

2025

- Answered domain-specific queries with **87%** relevance accuracy by building a Retrieval-Augmented Generation pipeline using LangChain, FAISS, and a locally hosted LLM.
- Indexed **500+** document chunks via sentence-transformer embeddings and semantic similarity retrieval to ground responses in source material.
- Exposed the pipeline through a FastAPI endpoint supporting streamed HTTP responses, reducing average query response time to under **2 seconds**.

Cloud Infrastructure Automation: AWS and Terraform

2023

- Provisioned **4** repeatable AWS environments (EC2, VPC, S3, IAM) using Terraform modules, cutting setup time from 3 hours to under **10 minutes**.
- Deployed containerised applications on Kubernetes with automated Docker builds, achieving zero-downtime releases via a GitHub Actions delivery pipeline.

EDUCATION

Gogte Institute of Technology

2021 – 2024

Bachelor of Engineering, Information Science and Engineering

CERTIFICATIONS

Google Cybersecurity Professional Certificate, Google – 2024

Introduction to Cybersecurity, Cisco – 2024